

CarpetX

2021 Einstein Toolkit Summer School, UIUC

Erik Schnetter, Perimeter Institute

2021-07-27



Flesh and Thorns

Plug-In "Thorns" (modules)

extensible APIs

remote steering

ANSI C

driver

parameters

Fortran/C/C++

input/output

scheduling

interpolation

Core "Flesh"

equations of state

error handling

SOR solver

make system

Your Physics !!

wave evolvers

grid variables

**Your Computational
Tools !!**

multigrid

utilities

coordinates

boundary conditions

black holes

The Cactus framework


- The **Driver** thorn in Cactus is the “main” function of the simulation:
- Scheduling (calls physics functions)
- Parallelism (splits grid functions across processes, exchanges ghost zones)
- AMR (levels, boxes, interpolation) (**AMReX!**)
- I/O (reading, writing, checkpointing)
- ...

Outline

- Tasks of a driver thorn
- Why AMR (Adaptive Mesh Refinement)
- CarpetX: New features (and missing features)
- New safety features
- Vertex and cell centering (conservation laws, div B)
- Parallelisation and efficiency (MPI, OpenMP, GPUs)
- I/O, SIMD, ...


Adaptive Mesh Refinement

- Need very high resolution only in small part of domain
- Often used for **compressible hydrodynamics**:
 - Moving shocks, discontinuities, surfaces



[AMReX documentation]

[Wikipedia: AMR]




DB: gaussian.it000000000.silo

Cycle: 0

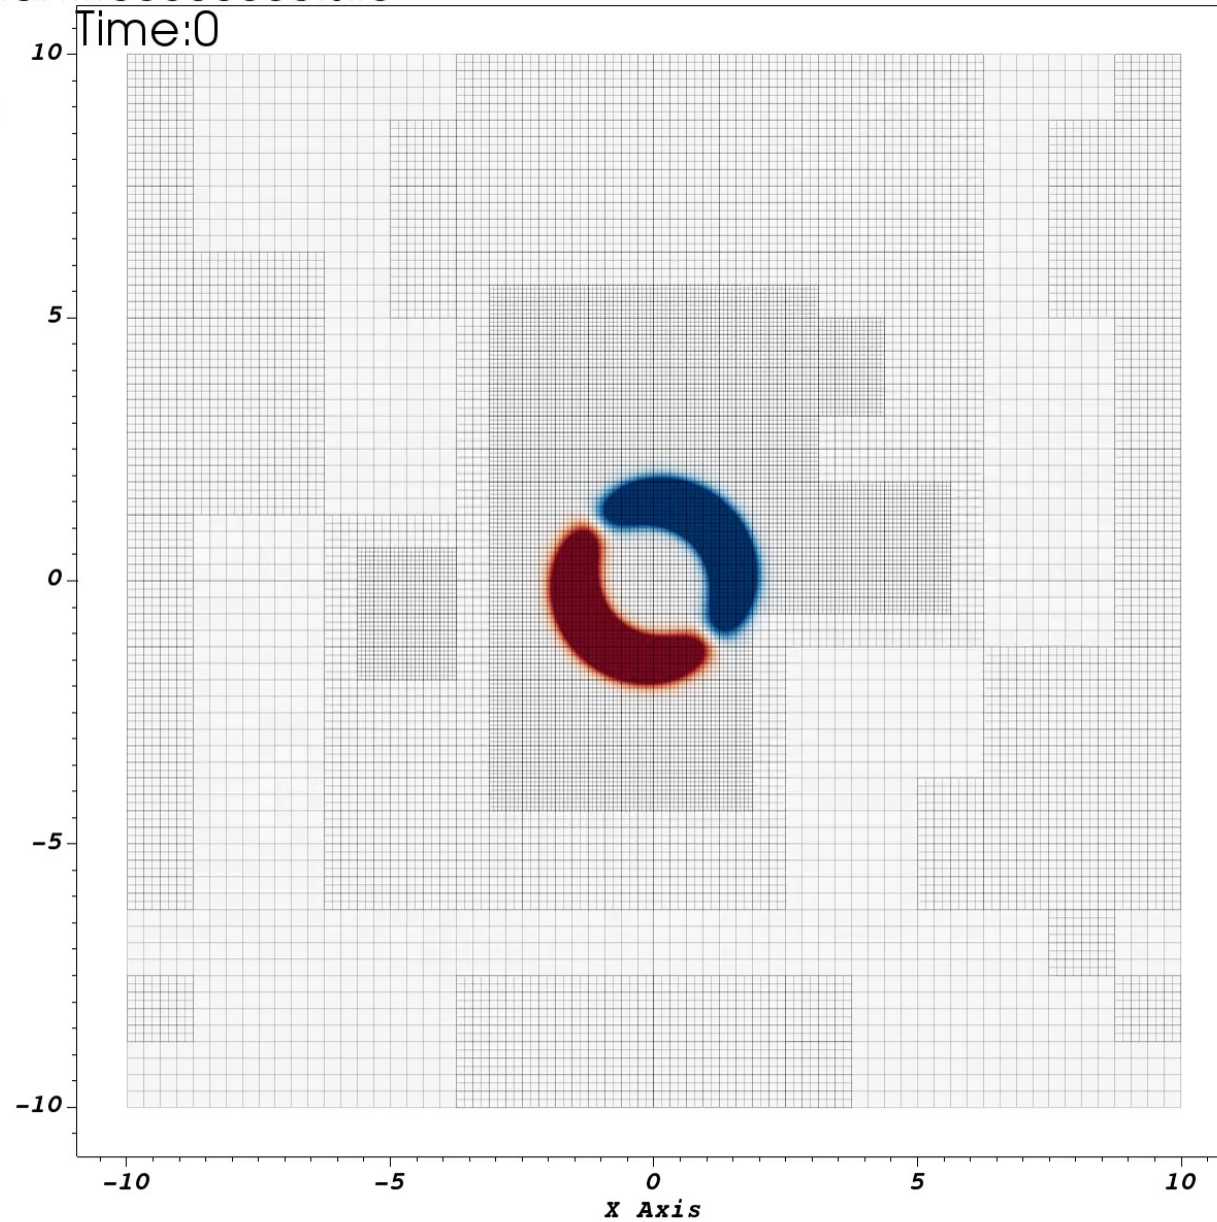
Time:0

Pseudocolor
Var: kleingordon_phi



Mesh
Var: gh

Y Axis




Scalar wave on Minkowski background

Lucas Sanches

Universidade Federal do ABC (Brazil)

CarpetX – a new mesh refinement driver for numeric relativity



- Designed for the Einstein Toolkit
- True adaptive mesh refinement based on local error estimate
- High-order prolongation/ restriction operators
- Vertex/cell/face/edge-centred variables
- Refluxing for exact conservation in (M)HD
- can currently run a qc0 BBH merger
- see presentations by Don Willcox and Erik Schnetter
- Uses AMReX
- for Exascale scalability
- Much improved multi-threading
- Much improved I/O speed (ADIOS2, openPMD)
- Improved SIMD vectorization
- Works with GPUs (CUDA)

CarpetX: Missing features

- Not yet implemented:
 - Global time stepping; no subcycling in time
 - Rotational symmetry boundaries
 - GRMHD code
 - NRPy+ backend
- Currently unknown (untested):
 - Recipes for efficient GPU kernels
 - MPI scalability (but AMReX is scalable)

New Safety Features

Scheduling

- Initial conditions:

1. CCTK_BASEGRID
2. Loop:
 1. Initialisation done?
 2. Regrid
 - CCTK_BASEGRID
 - CCTK_POSTREGRID
 3. CCTK_INITIAL
 4. CCTK_POSTINITIAL
3. CCTK_POSTSTEP
4. CCTK_ANALYSIS
5. OutputGH

- Evolution:

1. Evolution done?
2. Regrid
 - CCTK_BASEGRID
 - CCTK_POSTREGRID
3. CCTK_PRESTEP
4. CCTK_EVOL
5. CCTK_POSTSTEP
6. CCTK_ANALYSIS
7. OutputGH

Scheduling

- BASEGRID:
 - Set up constant data (e.g. coordinates)
- INITIAL, POSTINITIAL:
 - Initialise state vector (including boundaries)
 - Define local error for regridding
- EVOL:
 - Evolve state vector (and set boundaries)
- POSTSTEP, ANALYSIS:
 - Calculate ephemeral values (e.g. constraints)
 - Define local error for regridding
- POSTREGRID:
 - Re-apply boundary conditions to state vector

schedule.ccl:

```
SCHEDULE GROUP WaveToyCPU_RHSGroup IN ODESolvers_RHS
{
} "Calculate RHS"
```

```
SCHEDULE WaveToyCPU_RHS IN WaveToyCPU_RHSGroup
{
  LANG: C
  READS: phi(everywhere) psi(everywhere)
  WRITES: phirhs(interior) psirhs(interior)
} "Calculate RHS for the wave equation"
```

```
SCHEDULE WaveToyCPU_RHSSync IN WaveToyCPU_RHSGroup AFTER WaveToyCPU_RHS
{
  LANG: C
  SYNC: rhs
} "Boundary conditions for the RHS of the wave equation"
```

```
SCHEDULE WaveToyCPU_RHSBoundaries IN WaveToyCPU_RHSGroup AFTER WaveToyCPU_RHSSync
{
  LANG: C
  WRITES: phirhs(boundary) psirhs(boundary)
} "Boundary conditions for the RHS of the wave equation"
```

Declaring Dependencies in the Schedule

- Each scheduled function needs to declare which parts of which variables it reads or writes
- Regridding, synchronization, prolongation etc. do this implicitly as well
- The driver checks consistency, and flags errors
- The driver cross-checks these declarations via:
 - Undefined variables
 - Const variables
 - Variables set to nan, and checked for nan
 - Checksums of parts of variables
- Also, there are “informative” error messages

QC-0 (black hole binary)

Large box (+/- 128)


$dx[\text{coarse}] = 1$

$dx[\text{fine}] = 1/24$

4th order everything

~3/4 orbit before merging


DB: qc0.it000000000.silo
Cycle: 0 Time:0



Vertex and Cell centering

Vertex and Cell centering

- "Points" can be located at vertices, edges, faces, and interior of cells (in 3D)
 - Sample function at vertex
 - Integrate function along edge
 - Average function over cell
- Allows "baking in" certain properties into discretization scheme
 - **topological invariants**
 - E.g. $\text{div } \mathbf{B} = 0$, baryon number conservation



interface.ccl:

```
CCTK_REAL state TYPE=gf TAGS='index={1 1 1} rhs="rhs"'
{
  phi psi
} "Scalar potential for wave equation"
```

```
CCTK_REAL rhs TYPE=gf TAGS='index={1 1 1} checkpoint="no"'
{
  phirhs psirhs
} "RHS for scalar potential for wave equation"
```

```
CCTK_REAL energy TYPE=gf TAGS='index={1 1 1} checkpoint="no"'
{
  eps
} "Energy density for wave equation"
```

```
CCTK_REAL err TYPE=gf TAGS='index={1 1 1} checkpoint="no"'
{
  phierr psierr
} "Error of wave equation solution"
```


Parallelisation and Efficiency

Parallelisation and Efficiency

- **Shared** memory: Many cores can access the same memory
 - Need to split loops, but can use same array
 - Easy to use, but only for small problems (workstation)
 - Multi-**threading** (OpenMP)
- **Distributed** memory: No common memory
 - Need to split data structures as well as loops
 - Difficult but necessary (HPC systems)
 - Separate **processes** (MPI)
- **Accelerators** (“GPUs”):
 - Each accelerator is a shared memory system
 - HPC systems have many accelerators, i.e. distributed memory

Ghost Zones

- **Tiling** for shared memory
 - Each block has 4 coloured tiles
- **Ghost zones** for distributed memory
 - Each block is surrounded by a layer of extra points
 - Ghost zones need to be kept consistent (**synchronization**)
- AMR:
 - **Prolongation** (interpolation from coarse to finer level ghosts)
 - **Restriction** (from fine to coarser level where they overlap)




DB: gaussian.it000000000.silo

Cycle: 0

Time:0


Pseudocolor
Var: kleingordon_phi



Max: 0.4595
Min: -0.4595

Mesh
Var: gh

Y Axis



Scalar wave on Minkowski background

Lucas Sanches

Universidade Federal do ABC (Brazil)

Miscellaneous

I/O

- Need standardized file formats and standardized metadata
 - For post-processing, visualization, initial data exchange, etc.
- (Low-level) file formats can store blocks of data
 - HDF5, ASDF, ADIOS2, ...
- (High-level) metadata describes how blocks need to be assembled
 - Silo, openPMD, ???
- Efficient parallel I/O is difficult
 - **ADIOS2**, with **openPMD** metadata

SIMD (vectorization for CPUs)

- Modern CPUs can execute up to 8 operations simultaneously
 - (“threads” in CUDA, CPU threads are “blocks” in CUDA)
- Compilers are not good at generating SIMD code from scalar kernels